

AI Project: Predicting Patient Readmission Risk

Author : ODUGBILE OLALEKAN AYODELE

Part 1: Short Answer Questions

1. Problem Definition

Hypothetical AI Problem: Predicting student dropout rates in an online university.

Objectives:

- To identify at-risk students early in the semester.
- To enable targeted interventions such as academic counseling or resource allocation.
- To understand the key factors contributing to student dropout.

Stakeholders:

- Students: The primary beneficiaries of improved support.
- University Administration: Responsible for student retention and resource management.

Key Performance Indicator (KPI): Precision of the model. This is crucial because we want to minimize false positives—ensuring that interventions are directed at students who are genuinely at risk, rather than wasting resources on those who are not.

2. Data Collection & Preprocessing

Data Sources:

Learning Management System (LMS) Data: Login frequency, assignment submission times, quiz scores, and forum participation.

Student Information System (SIS) Data: Demographic information, past academic performance, and course enrollment history.

Potential Bias: Socioeconomic Bias. If the model is trained primarily on data from students with reliable internet access and personal computers, it may perform poorly for students who access courses primarily via mobile devices or public computers, potentially misclassifying them as at-risk due to lower engagement metrics.

Preprocessing Steps:

Handling Missing Data: Impute numerical features (e.g., average quiz score) with the median and categorical features (e.g., major) with the mode. If a significant portion of data is missing, the feature may be dropped.

Normalization/Standardization: Scale numerical features (e.g., number of logins, time spent) to a common range (like 0-1) or standard distribution to ensure no single feature dominates the model due to its scale.

Encoding Categorical Variables: Convert categorical data (e.g., department, country) into numerical format using one-hot encoding to make it interpretable for the model.

3. Model Development

Chosen Model: Gradient Boosting Machine (e.g., XGBoost).

Justification: It provides high accuracy, handles mixed data types (numerical and categorical) well, and can model complex, non-linear relationships between student activities and dropout risk. It also provides feature importance, which helps in understanding the key factors leading to dropout.

Data Splitting:

Training Set (70%): Used to train the model.

Validation Set (15%): Used for hyperparameter tuning and model selection during development.

Test Set (15%): Used only once, at the very end, to provide an unbiased evaluation of the final model's performance.

Hyperparameters to Tune:

`learning_rate`: Controls how quickly the model adapts to the problem. A lower rate is more robust but requires more trees.

`max_depth`: Controls the complexity of each individual tree. Tuning this helps balance model performance and prevents overfitting.

4. Evaluation & Deployment

Evaluation Metrics:

F1-Score: This is the harmonic mean of precision and recall. It is ideal for this imbalanced classification problem where both false positives (wasting resources) and false negatives (missing at-risk students) are important.

AUC-ROC (Area Under the Receiver Operating Characteristic Curve): Measures the model's ability to distinguish between the classes (will drop out vs. will not) across all classification thresholds.

Concept Drift & Monitoring:

Concept Drift occurs when the statistical properties of the target variable (what we are predicting) change over time. For example, a global event or a new curriculum could alter the factors that influence student dropout.

Monitoring: I would monitor it by tracking the model's performance metrics (F1-score, accuracy) on a held-out recent sample of data over time. A consistent drop in performance would signal potential concept drift.

Technical Challenge in Deployment:

Scalability & Latency: The system must generate predictions quickly for thousands of students without slowing down the university's main systems. This can be addressed by

deploying the model as a scalable microservice using cloud services (e.g., AWS SageMaker, Google AI Platform) and implementing efficient data pipelines.

Part 2: Case Study Application

Problem Scope

Problem: Develop a binary classification system to predict the risk of a patient being readmitted to the hospital within 30 days of discharge.

Objectives:

Reduce preventable readmissions, improving patient outcomes.

Enable care teams to provide proactive post-discharge care (e.g., follow-up calls, extra support).

Optimize hospital resource allocation and reduce healthcare costs.

Stakeholders: Patients, Physicians/Nurses, Hospital Administrators, and Insurance Payers.

Data Strategy

Data Sources:

Electronic Health Records (EHRs): Diagnosis codes, medications prescribed, length of stay, procedures performed, lab results.

Demographic & Socioeconomic Data: Age, gender, race, and potentially area-level data on income and education (to be used with extreme caution due to bias risks).

Ethical Concerns:

Patient Privacy: Handling of Protected Health Information (PHI) is governed by strict regulations like HIPAA. Data must be anonymized or de-identified for training, and access must be rigorously controlled.

Algorithmic Bias: The model could learn and amplify historical biases in healthcare delivery. For example, if certain demographic groups have historically received less effective care, they might have higher readmission rates, and the model could unfairly target them as "high risk" without addressing the root cause.

Preprocessing & Feature Engineering Pipeline:

Data Cleaning: Handle missing lab values (impute with normal range median), remove duplicate records.

Feature Engineering:

- Create a `number_of_previous_admissions` feature.
- Create comorbidity scores (e.g., Charlson Comorbidity Index) from diagnosis codes.
- Create a polypharmacy flag for patients on a large number of medications.
- Encoding: One-hot encode categorical variables like `primary_diagnosis`.
- Normalization: Standardize numerical features like `age` and `length_of_stay`.

Model Development (10 points)

Chosen Model: Logistic Regression.

Justification: In a high-stakes field like healthcare, model interpretability is often more critical than marginal gains in accuracy. Logistic Regression provides easily explainable

coefficients, showing how each feature (e.g., a specific medication) contributes to the readmission risk. This allows doctors to trust and understand the model's predictions.

Confusion Matrix & Calculations (Hypothetical Data for 100 patients):

- Actual Readmitted: 15
- Actual Not Readmitted: 85
- The model predicts:
- True Positives (TP): 10
- False Positives (FP): 8
- True Negatives (TN): 77
- False Negatives (FN): 5

Precision = $TP / (TP + FP) = 10 / (10 + 8) = 0.556$ (When the model says "high risk," it is correct 55.6% of the time).

Recall = $TP / (TP + FN) = 10 / (10 + 5) = 0.667$ (The model identifies 66.7% of all patients who were actually readmitted).

Deployment (10 points)

Integration Steps:

API Development: Package the trained model as a REST API.

HIPAA-Compliant Cloud: Deploy the API on a HIPAA-compliant cloud server (e.g., AWS with a BAA).

EHR Integration: The hospital's EHR system would send anonymized patient data at the time of discharge to this API.

Dashboard Display: The risk score is returned and displayed in a secure dashboard for the care team, flagging high-risk patients.

Ensuring Regulatory Compliance (HIPAA):

Data De-identification: All patient data used for model inference must be stripped of direct identifiers.

Access Controls: Implement strict role-based access control to the model and its predictions.

Business Associate Agreement (BAA): Sign a BAA with any third-party cloud provider to ensure they are legally obligated to protect PHI.

Audit Trails: Maintain logs of who accessed the model and which predictions were made.

Optimization

Method to Address Overfitting: L1 Regularization (Lasso).

Explanation: L1 regularization adds a penalty to the loss function equal to the absolute value of the magnitude of coefficients. This forces the model to shrink the coefficients for less important features to zero, effectively performing feature selection and producing a simpler, more generalizable model that is less likely to overfit to the noise in the training data.

Part 3: Critical Thinking

Ethics & Bias

Impact of Biased Data: Biased training data could lead to systematically underestimating the readmission risk for certain patient groups (e.g., racial minorities). This would create a "blind spot" for the care team, denying these patients the proactive support they need and potentially worsening health disparities. Conversely, over-predicting risk for a group could lead to intervention fatigue and the unnecessary allocation of limited resources.

Mitigation Strategy: Pre-processing Fairness Audit. Actively analyze the training data for representation disparities across different demographic groups. Techniques like re-sampling (over-sampling the underrepresented group) or re-weighting (assigning higher importance to instances from underrepresented groups) can be applied to create a more balanced dataset before the model is even trained.

Trade-offs (10 points)

Interpretability vs. Accuracy: In healthcare, a slightly less accurate but interpretable model (like Logistic Regression) is often preferred over a "black box" model (like a complex Neural Network). Doctors need to understand the "why" behind a prediction to trust it and make informed clinical decisions. A highly accurate model is useless if medical professionals cannot interpret its reasoning and integrate it into their care plan.

Impact of Limited Computational Resources: Limited resources would strongly favor simpler, more efficient models like Logistic Regression or Decision Trees over computationally intensive ones like large Neural Networks or ensemble methods. The focus would shift from maximizing predictive performance to finding the most effective model that can run within the hardware constraints, potentially on-premise without expensive cloud computing.

Part 4: Reflection & Workflow Diagram (10 points)

Reflection (5 points)

Most Challenging Part: The most challenging part was navigating the trade-offs in the Case Study, specifically balancing model performance with ethical and practical constraints. Choosing a model for the hospital required prioritizing interpretability and fairness over pure accuracy, which is a complex, non-technical decision. Ensuring HIPAA compliance added another layer of complexity that goes beyond pure machine learning.

Improvements with More Time/Resources:

Advanced Explainable AI (XAI): I would use tools like SHAP or LIME with a more complex model (like XGBoost) to provide both high accuracy and local explanations for each prediction, offering the "best of both worlds."

Robust Fairness Testing: Implement a more rigorous, continuous fairness evaluation framework to monitor the model for disparate impact across different subgroups in production.

Pilot Deployment & A/B Testing: Run a controlled clinical trial to measure the real-world impact of the AI system on readmission rates and patient outcomes before a full-scale rollout.

AI Development Workflow



